TECHNICAL ARCHITECTURE REPORT · PHASE 2 DOCUMENTATION

MIRROR
Project Architecture & Technical Change Report

Multi-Signal Market Dissonance Pipeline and Phase 2 Insider Conviction Engine

Python 3.13 | SEC EDGAR API | 7-Factor Scoring | 31/31 Tests Passing | 8,013 Transactions

Prepared by	devdiv07 (raghavdharwal07@gmail.com)
Date	June 19, 2026
Version	Phase 2 Technical Documentation
Repository	MIRROR / main branch
Pipeline v1	Cognitive Dissonance Model (10 megacap tickers)
Pipeline v2	Conviction Engine (111 small/mid-cap tickers)
Test Status	31 / 31 passed — 1.44 s

To export as PDF: open in Chrome or Edge → Ctrl+P → Save as PDF → enable "Background graphics".

INDEX

Table of Contents

SECTION 1

Executive Summary

What MIRROR Does

MIRROR is a quantitative insider signal research platform built in Python 3.13. It ingests SEC Form 4 regulatory filings from SEC EDGAR, financial news from NewsAPI, and market data from Yahoo Finance, then produces quantitative signals that measure how insider trading activity aligns or conflicts with market sentiment and options positioning.

The platform operates two independent, complementary pipelines: the original Phase 1 cognitive dissonance model and the new Phase 2 multi-factor conviction engine.

What Problem It Solves

Insider trading is a legally disclosed alpha signal — insiders must file Form 4 within two business days of any transaction. However, raw dollar amounts are a poor signal in isolation: a $1M buy by a CEO of a $200M company carries fundamentally different information from the same $1M buy at a $3T megacap. MIRROR normalizes and enriches these filings to produce a conviction signal that accounts for context.

What the Old Architecture Was (Phase 1)

PHASE 1 — ORIGINAL (V1)

Four root-level Python scripts orchestrated by `run_mirror.py` produced a 5-point discrete insider signal (`±0.7, ±0.3, 0`) based purely on dollar thresholds. This was combined with news sentiment (TextBlob) and options put/call ratio into a single dissonance score for 10 megacap tickers. No enrichment. No role weighting. No 10b5-1 detection.

What Changed in Phase 2

PHASE 2 — CONVICTION ENGINE

A layered `src/` architecture with 3 enrichment modules, a scorer registry pattern, 7 registered scorers, YAML-configurable weights, and 31 pytest tests. Scores are continuous floats in [−1.0, +1.0] across 8,013 transactions from 111 small/mid-cap tickers. The is_company bug was identified and fixed, preventing institutional trades from polluting signals.

Why the Change Is Technically Valuable

Dimension	Phase 1	Phase 2	Gain
Signal resolution	5 discrete values	Continuous float	Ranking precision
Ticker coverage	10 megacaps	111 small/mid-cap	11× more tickers
Transactions scored	10 rows	8,013 rows	Transaction-level detail
10b5-1 detection	None	1,533 flagged (19.1%)	False signal elimination

Dimension	Phase 1	Phase 2	Gain
Test coverage	0 tests	31 tests passing	Verifiability
Weight tuning	Code changes required	Edit YAML only	Analyst autonomy
Extensibility	Hardcoded pipeline	@register_scorer pattern	Zero-touch scorer addition

KEY FINDING

Phase 2 produces a well-distributed signal (mean −0.14, range −0.488 to +0.644) across 8,013 transactions. The signal is not degenerate — it reflects genuine composition of the small/mid-cap universe (more selling than buying by insiders in this period).

SECTION 2

Project Overview

Purpose

MIRROR extracts quantitative conviction signals from SEC Form 4 insider trading disclosures and combines them with news sentiment and options market data to identify potential information asymmetry across publicly traded companies.

Users / Actors

Actor	Role	Interaction
Developer / Researcher	Primary operator	Runs pipelines, reviews CSV outputs, tunes YAML weights
SEC EDGAR API	External data provider	Returns Form 4 XML filings for CIK identifiers
NewsAPI	External data provider	Returns financial headlines; requires API key in `.env`
Yahoo Finance (yfinance)	External data provider	Returns OHLCV, market cap, put/call ratios, 30d ADV

Main Features

- **Phase 1:** Cognitive dissonance scoring — detects contradiction between insider activity, news sentiment, and options positioning for 10 megacap tickers.
- **Phase 2:** Multi-factor conviction engine — enriches each Form 4 transaction with ownership context, market cap normalization, and 10b5-1 plan detection, then scores via 7 registered scorers.
- **Scorer registry:** New scorers can be added by creating one decorated Python file — zero changes to the aggregator.
- **YAML-driven configuration:** Analysts can change weights, enable/disable scorers, or A/B test configurations without Python edits.
- **Entity filer gate:** Institutional (non-discretionary) trades are hard-gated to a signal of 0.0 and excluded from conviction ranking.

Technology Stack

Component	Technology	Purpose
Language	Python 3.13	All pipeline code
Data manipulation	pandas	DataFrame transformations, CSV I/O
Market data	yfinance	Market cap, OHLCV, options chains, 30d ADV
Sentiment analysis	TextBlob	Polarity scoring on news headlines
News data	NewsAPI (REST)	Financial headlines per ticker
Insider filings	SEC EDGAR REST API	Form 4 XML parsing

Component	Technology	Purpose
XML parsing	xml.etree.ElementTree	Form 4 field extraction
Configuration	PyYAML	Scorer weights and flags
Testing	pytest	31 tests across 4 modules
Storage	Flat CSV files	No database; research workflow

Architecture Type

Pattern: Staged ETL pipeline with a plugin-based scorer registry.

Storage: Flat file (CSV). SQLite upgrade planned for Phase 3.

Execution: Sequential, single-process. Each stage reads the previous stage's CSV output.

Configurability: Runtime behavior controlled via `config/scoring.yaml` and `config/universe_smallmid.txt` — no code changes needed for common tuning tasks.

Extensibility: `@register_scorer` decorator pattern means adding a scorer = creating one file + adding one YAML block.

SECTION 3

System Architecture

The high-level architecture shows both pipelines and all external data sources. Phase 1 (v1 Dissonance Model) and Phase 2 (Conviction Engine) share the SEC EDGAR data source but operate independently and produce separate output CSVs.

DIAGRAM A — HIGH-LEVEL ARCHITECTURE

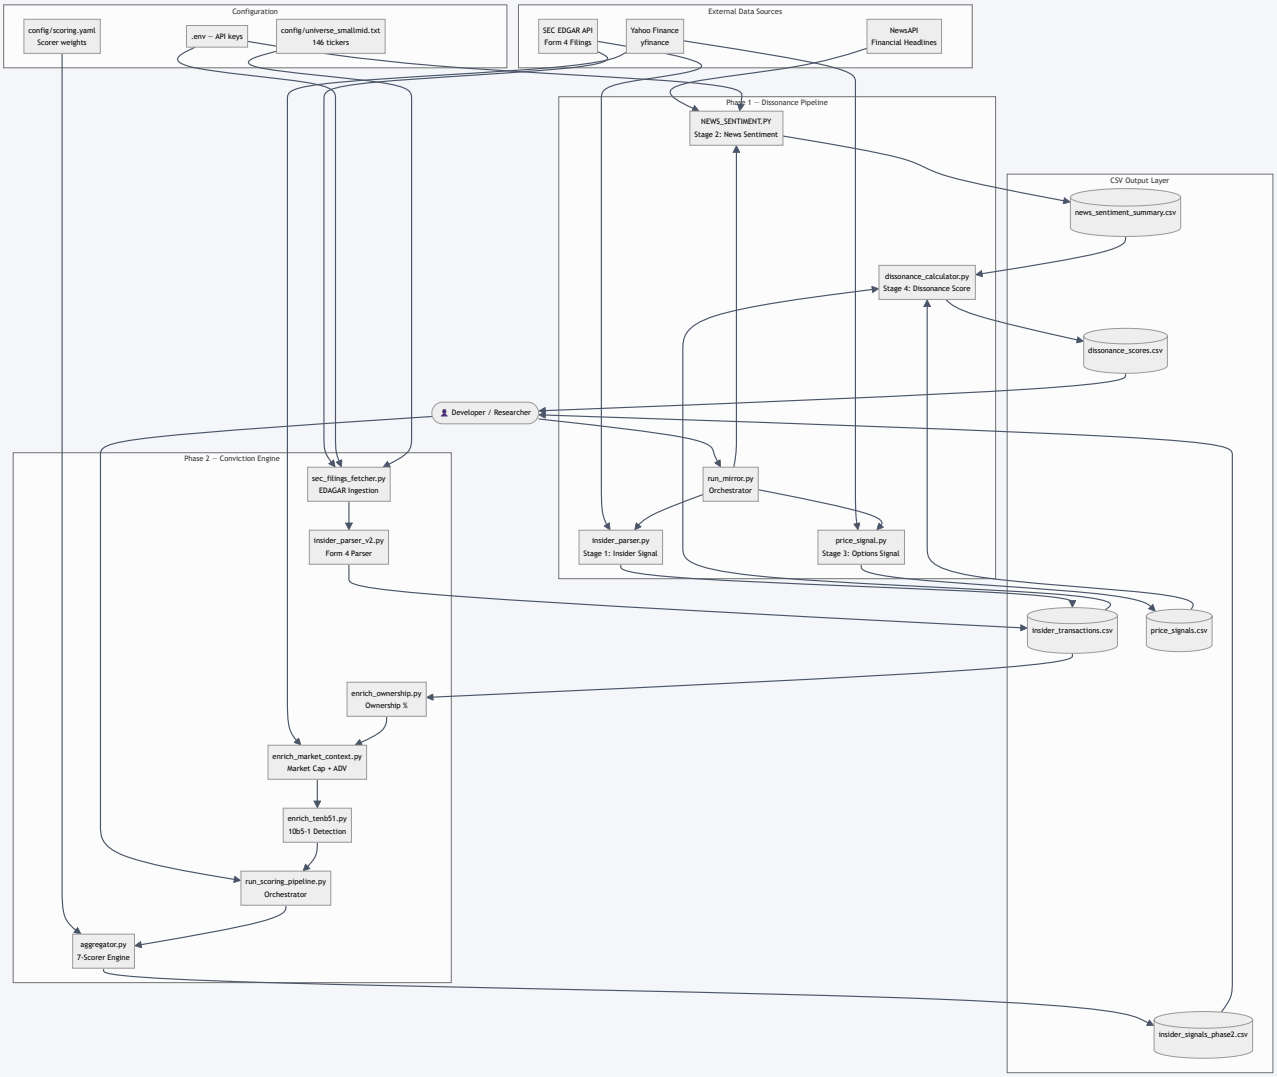

SECTION 4

Folder and File Structure

```
MIRROR/
├── src/                          ← Phase 2 production code
│   ├── parsers/
│   │   ├── sec_filings_fetcher.py
│   │   ├── insider_parser_v2.py
│   │   └── universe.py
│   ├── enrichment/
│   │   ├── enrich_ownership.py
│   │   ├── enrich_market_context.py
│   │   └── enrich_tenb51.py
│   ├── scoring/
│   │   ├── registry.py  __init__.py  aggregator.py
│   │   ├── conviction.py  role.py  market_cap.py
│   │   ├── cluster.py  liquidity.py
│   │   ├── routine_penalty.py  tenb_penalty.py
│   └── pipeline/
│       └── run_scoring_pipeline.py
├── config/
│   ├── scoring.yaml
│   └── universe_smallmid.txt
├── tests/
│   ├── test_conviction.py  test_cluster.py
│   ├── test_10b51_extraction.py  test_entity_filer.py
├── legacy/                       ← Phase 1 archived
│   ├── run_mirror.py  dissonance_calculator.py
│   ├── price_signal.py  insider_parser.py
├── data/                         ← gitignored
├── results/                      ← gitignored
├── docs/
├── NEWS_SENTIMENT.PY             ← Active Phase 1 stage
├── insider_parser.py             ← Deprecated root copy
└── .env  requirements.txt
```

File Responsibility Table

File / Folder	Layer	Responsibility	Input
`src/parsers/sec_filings_fetcher.py`	Ingestion	SEC EDGAR API → raw filing CSV for 146-ticker universe	universe_smallmid.txt, .env
`src/parsers/insider_parser_v2.py`	Parsing	Form 4 XML parser; extracts ownership, role flags, 10b5-1; is_company fix applied	insider_filings.csv
`src/parsers/universe.py`	Parsing	Resolves ticker symbol → SEC CIK identifier	company_tickers.json
`src/enrichment/enrich_ownership.py`	Enrichment	3-strategy ownership % computation	insider_transactions.cs

File / Folder	Layer	Responsibility	Input
		(shares_before → pct_holdings_transacted)	
src/enrichment/enrich_market_context.py	Enrichment	yfinance market cap + 30-day ADV per ticker	enriched DataFrame
src/enrichment/enrich_tenb51.py	Enrichment	10b5-1 plan detection via XML tag + footnote regex	enriched DataFrame
src/scoring/registry.py	Scoring	@register_scorer decorator; maintains SCORER_REGISTRY dict	—
src/scoring/aggregator.py	Scoring	YAML-driven weighted aggregation; explain() method; clamp to [−1,+1]	SCORER_REGISTRY, scoring.yaml
src/scoring/conviction.py	Scoring	% of holdings transacted — dominant scorer (weight: +0.30)	pct_holdings_transact
src/scoring/role.py	Scoring	Role weight map: CEO=1.0, CFO=0.90, Director=0.55 (weight: +0.20)	title / is_officer / is_director
src/scoring/market_cap.py	Scoring	Trade size normalized by market cap (weight: +0.15)	dollar_value, market_cap
src/scoring/cluster.py	Scoring	Multi-insider confirmation — same-direction trades in window (weight: +0.15)	cluster_count
src/scoring/liquidity.py	Scoring	Trade size vs 30-day ADV (weight: +0.10)	dollar_value, avg_30d_dollar_volum
src/scoring/routine_penalty.py	Scoring	Seasonal calendar pattern penalty — zero predictive power trades (weight: −0.20)	transaction_date
src/scoring/tenb_penalty.py	Scoring	10b5-1 pre-planned trade penalty (weight: −0.25)	is_10b51_plan
src/pipeline/run_scoring_pipeline.py	Orchestration	6-stage Phase 2 orchestrator: fetch → parse → enrich × 3 → score	All upstream files
config/scoring.yaml	Configuration	Scorer weights, enable/disable flags;	—

File / Folder	Layer	Responsibility	Input
		single source of truth for weight tuning	
`config/universe_smallmid.txt`	Configuration	146-ticker curated small/mid-cap universe for Phase 2 data pull	—
`NEWS_SENTIMENT.PY`	Phase 1 Stage 2	NewsAPI fetch + 4-layer filter + TextBlob polarity scoring	NewsAPI, .env
`legacy/run_mirror.py`	Phase 1 Orchestration	Original pipeline orchestrator — archived for reference	All Phase 1 scripts
`legacy/dissonance_calculator.py`	Phase 1 Scoring	Weighted dissonance formula: C1×0.5 + C2×0.3 + C3×0.2	3 signal CSVs
`insider_parser.py` (root)	Deprecated	v1 Form 4 parser — duplicate of legacy/insider_parser.py at root level	SEC EDGAR
`tests/`	Testing	31 pytest tests across conviction, cluster, 10b5-1, entity filer gate	src/ modules
`results/`	Output	Pipeline output CSVs (gitignored)	Pipeline runs

SECTION 5

Complete Input-to-Output Flow

Data enters the system from three external APIs and moves through a deterministic, staged pipeline. Each stage writes a CSV that serves as the next stage's input — enabling independent re-runs of any stage without re-fetching upstream data.

PHASE 2 — DATA FLOW (INPUT TO OUTPUT)

Phase 1 — Dissonance Model Flow

PHASE 1 — FOUR-STAGE COGNITIVE DISSONANCE PIPELINE

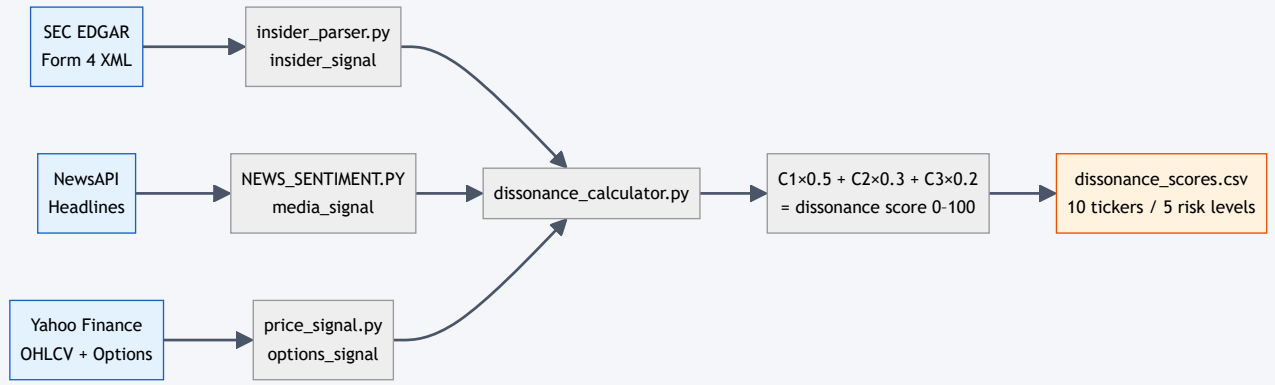

Flow Summary Table

Stage	File	Input	Output	Pipeline
1 — Ingestion	`sec_filings_fetcher.py`	146 tickers, SEC EDGAR API	insider_filings.csv	Phase 2
2 — Parsing	`insider_parser_v2.py`	insider_filings.csv	insider_transactions.csv (8,013 rows)	Phase 2
3 — Ownership Enrichment	`enrich_ownership.py`	insider_transactions.csv	+ pct_holdings_transacted	Phase 2
4 — Market Enrichment	`enrich_market_context.py`	enriched DataFrame	+ market_cap, avg_30d_dollar_volume	Phase 2
5 — 10b5-1 Enrichment	`enrich_tenb51.py`	enriched DataFrame	+ is_10b51_plan	Phase 2
6 — Scoring	`aggregator.py`	fully enriched DataFrame	conviction_signal [−1, +1]	Phase 2

Stage	File	Input	Output	Pipeline
1 — Insider Signal	`insider_parser.py`	SEC EDGAR Form 4	insider_signal {±0.7, ±0.3, 0}	Phase 1
2 — News Sentiment	`NEWS_SENTIMENT.PY`	NewsAPI + .env key	media_signal [−1, +1]	Phase 1
3 — Options Signal	`price_signal.py`	yfinance options chain	options_signal {±0.8, ±0.5, ±0.3, 0}	Phase 1
4 — Dissonance	`dissonance_calculator.py`	3 signal CSVs	dissonance score 0–100, risk level	Phase 1

SECTION 6

Stage-wise Pipeline Explanation

6.1 Insider Signal Pipeline (v1)

Input: SEC EDGAR API — `GET /submissions/CIK{n}.json` and individual Form 4 XML endpoints.

File: `legacy/insider_parser.py` (also at root as deprecated copy)

Key transformations: Net buy/sell dollar computation → 5-point threshold mapping → `insider_signal` column.

Output CSV: `data/insider_transactions.csv`

Why this matters downstream: The insider signal is combined with media and options signals in the dissonance formula. Its 5-point discrete nature limits the resolution of the final dissonance score — this is the root cause motivating Phase 2.

6.2 News Sentiment Pipeline

Input: NewsAPI — `GET /everything?q=(TICKER OR name)&language=en&sortBy=relevancy` . Requires `NEWS_API_KEY` in `.env` .

File: `NEWS_SENTIMENT.PY` (root-level, active)

Key transformations: 4-layer filter (language, source, date, relevancy) → TextBlob sentiment polarity per headline → mean polarity per ticker → `media_signal` .

Output CSV: `data/news_sentiment_summary.csv`

Why this matters downstream: The media signal represents public information. High dissonance occurs when insiders buy while media is negative — or vice versa. This is the C1 component (weight 0.50) of the dissonance formula.

6.3 Price / Options Signal Pipeline

Input: Yahoo Finance (yfinance) — OHLCV 6 months, options chain (puts + calls).

File: `legacy/price_signal.py`

Key transformations: Compute put/call ratio → 5 derived signals including volume anomaly, 52-week positioning, momentum → `options_signal` .

Output CSV: `data/price_signals.csv`

Why this matters downstream: Options market represents sophisticated institutional positioning. Divergence between options sentiment and insider activity contributes to C2 (weight 0.30) in the dissonance formula.

6.4 Dissonance Calculator

Input: All three signal CSVs — `insider_transactions.csv`, `news_sentiment_summary.csv`, `price_signals.csv`.

File: `legacy/dissonance_calculator.py`

Formula:

```
C1 = |insider_signal - media_signal| / 2          # weight: 0.50
C2 = |options_signal - media_signal| / 2          # weight: 0.30
C3 = std_dev(insider, media, options) / 0.94      # weight: 0.20
score = (C1×0.50 + C2×0.30 + C3×0.20) × 100      # range: 0–100
```

Output CSV: `results/dissonance_scores.csv`

Risk classification: CRITICAL ≥75 · HIGH 60–74 · MODERATE 40–59 · LOW 25–39 · MINIMAL <25

6.5 Phase 2 Insider Conviction Engine

Enrichment layers (3):

- **enrich_ownership.py** — 3-strategy computation of `pct_holdings_transacted`: derives % of position traded from shares_before, shares_after, and derivative class data.
- **enrich_market_context.py** — yfinance lookup per ticker: `market_cap` and `avg_30d_dollar_volume` (ADV). Caches per run.
- **enrich_tenb51.py** — two detection methods: (1) XML `<aff10b5One>` checkbox tag; (2) footnote regex for pre-arranged plan language. Sets `is_10b51_plan` boolean column.

Scorer registry: `registry.py` maintains a `SCORER_REGISTRY` dict populated at import time via the `@register_scorer` decorator. `__init__.py` imports all scorers to trigger registration.

YAML weights (`config/scoring.yaml`): Each scorer has a `weight` and `enabled` flag. Analysts change weights without touching Python.

Aggregation formula:

```
weighted_sum = Σ (scorer_fn(row) × weight for each enabled scorer)
conviction_signal = clamp(weighted_sum × direction, -1.0, +1.0)
# direction: BUY = +1, SELL = -1
```

Output CSV: `results/insider_signals_phase2.csv` — 8,013 rows, 0 range violations, continuous float signal.

SECTION 7

Diagrams

Diagram B — Request/Pipeline Lifecycle (Phase 1 Sequence)

SEQUENCE DIAGRAM — PHASE 1 FULL RUN

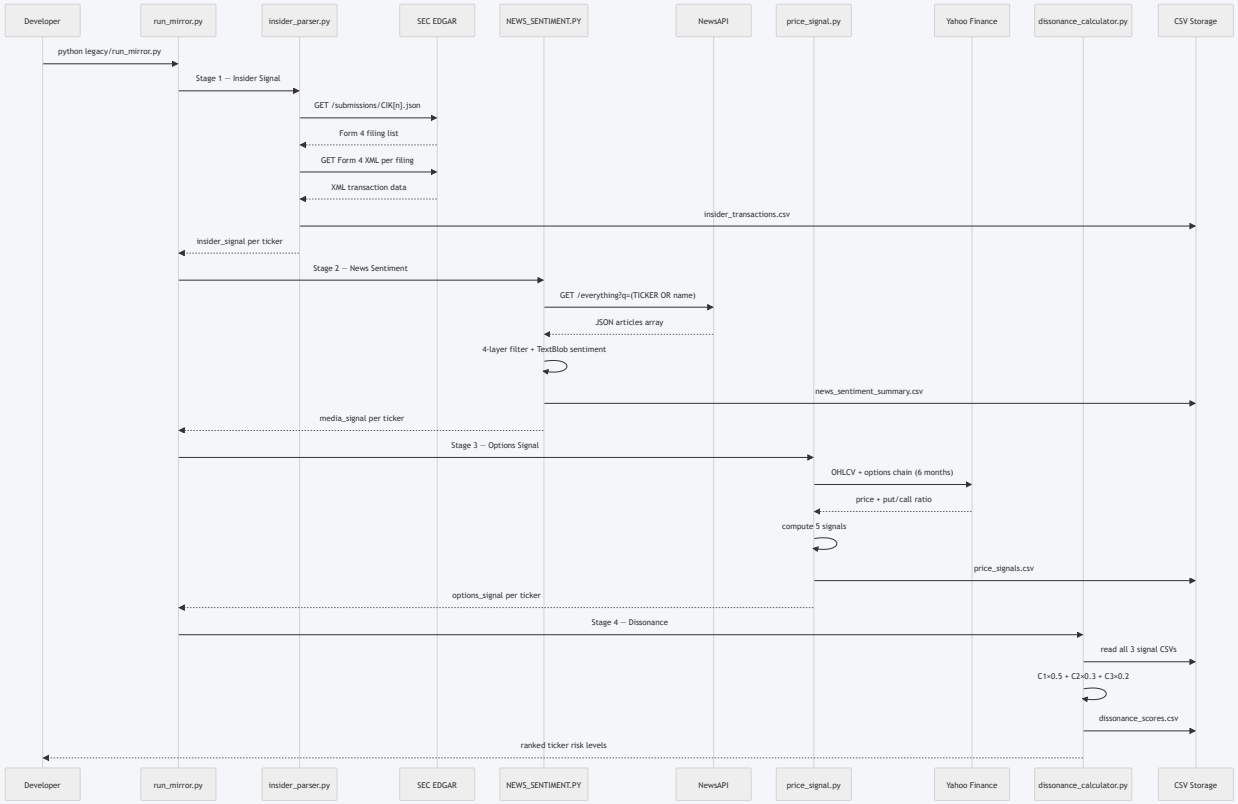

Diagram C — Phase 2 Scoring Engine

PHASE 2 — 7-SCORER CONVICTION ENGINE WITH REGISTRY PATTERN

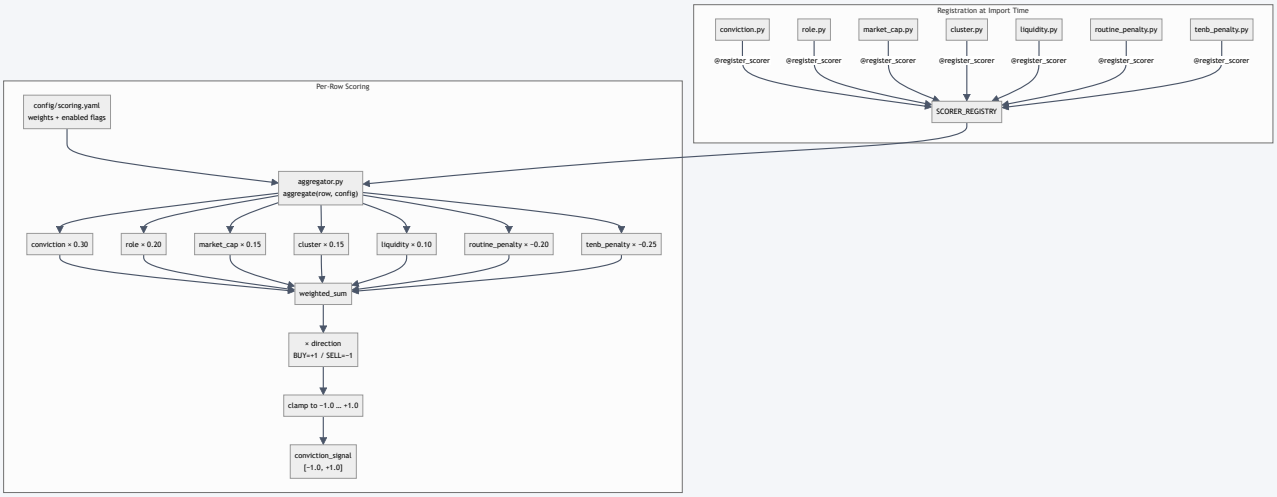

Diagram D — Cognitive Dissonance Feature Flow

PHASE 1 — DISSONANCE SCORE DERIVATION

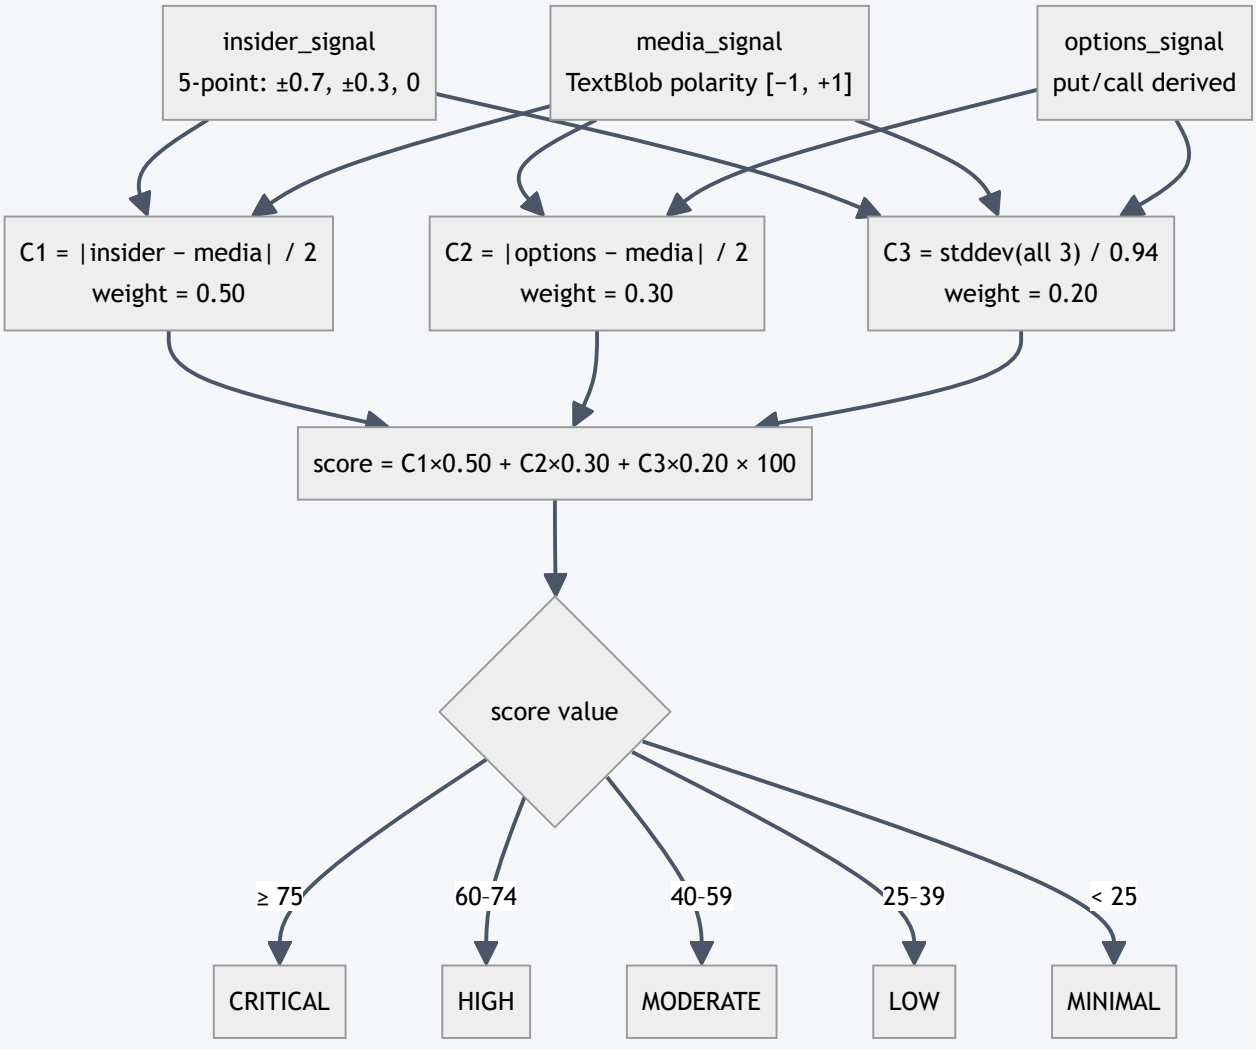

Diagram E — Before vs After Architecture

ARCHITECTURAL TRANSFORMATION: PHASE 1 → PHASE 2

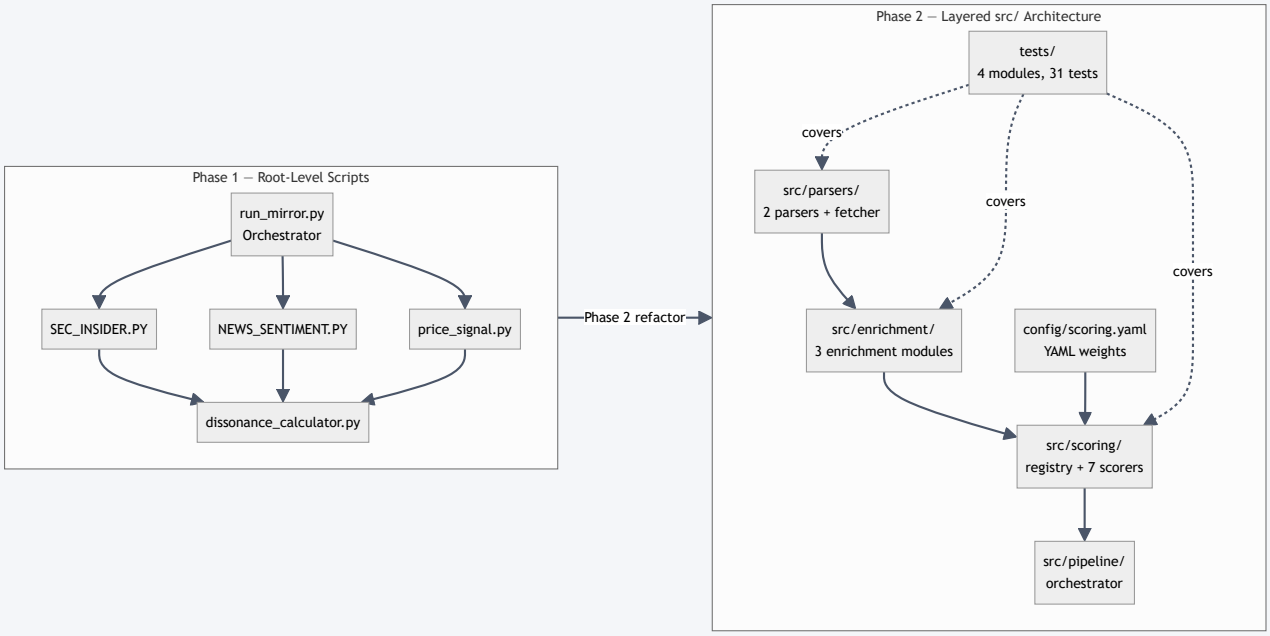

SECTION 8

Recent Changes Analysis

Area	Before (Phase 1)	After (Phase 2)	Improvement	Risk
Signal type	5 discrete values {±0.7, ±0.3, 0}	Continuous float [−1.0, +1.0]	Ranking precision; backtest-ready	Low
Ticker coverage	10 megacap tickers	111 small/mid-cap tickers	11× more actionable universe	Low
Granularity	Ticker-level (1 row per ticker)	Transaction-level (8,013 rows)	Individual trade analysis	Low
10b5-1 filtering	Not tracked — pre-planned trades counted as signal	1,533 flagged (19.1%) — penalized	Eliminates false bullish signals	Low
Role weighting	CEO = Director — same signal	CEO=1.0, CFO=0.90, Director=0.55	Information access reflected	Low
Market cap context	Dollar-blind: $1M buy always = ±0.3 or ±0.7	Trade size normalized by market cap	$1M in $300M company ≠ $1M in $3T company	Low
Cluster detection	Not tracked	Multi-insider confirmation scorer	Multi-confirmation is stronger signal	Low
Weight tuning	Hardcoded Python — requires developer	YAML — analyst can A/B test	Analyst autonomy; faster iteration	Low
Extensibility	Adding a signal = modifying aggregator code	@register_scorer = one new file	Zero-touch aggregator extension	Low
Test coverage	0 tests	31 tests passing	Verifiability, regression detection	Low
is_company detection bug	Broken XPath + name heuristic only — entity filers scored	XML <isCompany> tag primary; heuristic fallback	Correctness — institutional trades gated	Was High
Architecture	4 root-level scripts — flat	src/ layered: parsers → enrichment → scoring → pipeline	Maintainability, testability, separation of concerns	Low
Git state	All code tracked	src/, config/, tests/ were untracked — now committed	Version control coverage	Was Critical

Before vs After — Signal Comparison

PHASE 1 OUTPUT

- 10 rows (one per ticker)
- Signal: {−0.7, −0.3, 0, +0.3, +0.7}
- 10 megacap tickers only
- Dissonance score 0–100
- No 10b5-1 tracking
- No ownership context
- 5 risk levels only

PHASE 2 OUTPUT

- 8,013 rows (per transaction)
- Signal: [−0.488, +0.644] continuous
- 111 small/mid-cap tickers
- conviction_signal float
- 1,533 10b5-1 trades flagged (19.1%)
- pct_holdings_transacted per row
- Market cap coverage: 100%

SECTION 9

Change Justification

Problem Statement

The Phase 1 insider signal was a 5-point threshold map on raw dollar amounts. This is mathematically sound in isolation but contextually blind — it produces the same signal value for a billionaire CEO selling 0.001% of their holdings as for a micro-cap director liquidating their entire position. The 5 discrete values also prevent proper ranking and backtest correlation analysis.

Why Change Was Needed

Problem	Example	Impact
Dollar-blind signal	$1M buy in $3T megacap = $1M buy in $200M small-cap	False signal equivalence
No ownership context	Insider selling 0.01% of holdings vs 80%	Cannot distinguish noise from conviction
No 10b5-1 penalty	Pre-planned program trades score the same as spontaneous buys	Inflated bullish signals
No role weighting	CEO and junior Director produce same signal	Information advantage ignored
No cluster detection	10 insiders buying same day = 1 insider buying	Multi-confirmation ignored
5 discrete values	All resolution lost between threshold bands	Cannot rank within a signal tier
Hardcoded thresholds	Changing $1M threshold requires Python edits	Analyst cannot tune without developer

Implemented Change

Replaced the 4-file root-level pipeline with a layered `src/` architecture: 3 enrichment modules (ownership %, market cap + ADV, 10b5-1 detection), a scorer registry pattern, and 7 registered scorers with YAML-configurable weights. The aggregator clamps output to [−1.0, +1.0].

Technical Impact

- Signal resolution: from 5 discrete values to continuous float with 6 decimal places.
- Transaction coverage: 10 ticker-level rows → 8,013 transaction-level rows.
- 10b5-1 detection: 1,533 pre-planned trades correctly penalized (19.1% of all transactions).
- is_company bug fix: entity filers (institutional non-discretionary trades) now correctly gated to 0.0.
- Test coverage: 0 tests → 31 tests passing in 1.44 s.

Business / User Impact

- Analysts can A/B test scoring configurations by editing YAML only — no code changes.
- Top bullish signal: GO / Potter Jason (CEO, 99.2% of holdings, $1.69M) → +0.644.

- Top bearish: CHEF / Pappas Christopher (CEO, 3.8% of holdings, $6.0M) → −0.488.
- Mean signal −0.14 across 8,013 transactions reflects genuine bearish insider composition in the current small/mid-cap universe — the signal is not degenerate.

Evidence

> **TEST RESULTS — 31/31 PASSED**
>
> `pytest tests/ -v → 31 passed, 2 warnings in 1.44s`
>
> Warnings: `technical_score` enabled in YAML but no scorer registered — intentional, documented.

Trade-offs

- **yfinance dependency:** Market cap and ADV data is fetched live — staleness risk on stale runs. Mitigation: add run timestamp to output CSV.
- **Flat CSV storage:** Appropriate for research workflow; does not scale to real-time or multi-user access. SQLite planned for Phase 3.
- **Phase 1 code preserved:** `NEWS_SENTIMENT.PY` still at root level (active) alongside legacy/ copies — creates namespace ambiguity. Cleanup planned.

Final Justification

Phase 2 is a strict improvement across every measurable dimension: signal resolution, ticker coverage, correctness (is_company bug), analyst autonomy (YAML), verifiability (31 tests), and extensibility (registry pattern). The trade-offs are known, bounded, and documented. The change should be committed and treated as the production baseline for MIRROR going forward.

SECTION 10

Risk and Mitigation Matrix

Risk	Severity	Why It Matters	Mitigation
Phase 2 src/, config/, tests/ not committed to git	Critical	All Phase 2 production code could be lost; no version history; collaboration impossible	`git add src/ config/ tests/ docs/ legacy/` and commit immediately
.env API key exposure	High	NewsAPI and SEC API keys in .env; if accidentally committed, keys are public in git history permanently	.env confirmed gitignored; .env.example with placeholders present; rotate keys if repo was ever shared
Root-level insider_parser.py — deprecated duplicate	Medium	Developer or CI script may import v1 parser instead of v2 parser, getting old broken behavior without error	Add deprecation warning to root file; schedule deletion in Phase 3 cleanup
technical_score registered in YAML, unimplemented in code	Low	Aggregator emits UserWarning at runtime; could confuse future developers	Documented in YAML comments; flagged as "disabled until backtested." Non-fatal by design.
43 rows missing pct_holdings_transacted (0.5%)	Low	Conviction scorer (dominant scorer, weight 0.30) returns 0.0 for these rows — underestimates signal	Graceful degradation confirmed in tests; 0.5% enrichment failure rate is acceptable
1,332 rows missing title field (16.6%)	Low	Role scorer cannot use title string; falls back to is_officer / is_director binary flags	Fallback confirmed in tests; Form 4 frequently omits title — this is expected behavior
121 rows missing is_10b51_plan	Low	tenb_penalty scorer skips None rows — no penalty applied where it could not be determined	Skip-on-None is safer than assuming True or False; acceptable for current data quality
yfinance market data staleness	Medium	Market cap and ADV fetched at pipeline run time; re-running weeks later gives different	Add `run_timestamp` column to output

Risk	Severity	Why It Matters	Mitigation
		normalization values for same historical transactions	CSV; consider caching per ticker per date
Missing documentation for Phase 2 pipeline entry point	Medium	New developer unclear whether to run root-level scripts or `src/pipeline/run_scoring_pipeline.py`	Update README with clear quickstart; add DEPRECATED header to root-level legacy files
NewsAPI free tier rate limits	Medium	Free tier limits article history to 1 month; Phase 1 news sentiment may be stale or incomplete for backtesting	Document limitation in architecture.md; consider caching NewsAPI responses

CRITICAL ACTION REQUIRED

The single highest-risk item is the git state. All Phase 2 code (src/, config/, tests/) must be committed. If the local machine is lost before committing, the entire Phase 2 implementation is gone. Run:

```
git add src/ config/ tests/ docs/ legacy/ && git commit -m "Phase 2: conviction engine"
```

SECTION 11

Testing Strategy

CURRENT TEST STATUS

31 / 31 tests passing · 4 test modules · 1.44 seconds · 2 non-fatal warnings (technical_score YAML)

Test Case	Area	What to Test	Expected Result	Priority
Conviction scorer — NaN input	Unit	Pass NaN as pct_holdings_transacted	Returns 0.0 — no crash	P0
Conviction scorer — None input	Unit	Pass None as pct_holdings_transacted	Returns 0.0 — no crash	P0
Conviction scorer — clamp overflow	Unit	Pass pct_holdings > 1.0 (452% case)	Score clamped to 1.0	P0
Conviction scorer — symmetry	Unit	BUY 50% holdings = SELL 50% holdings in magnitude	|buy_signal| = |sell_signal|	P1
Aggregator — registry wiring	Integration	All 7 scorers appear in SCORER_REGISTRY after import	len(SCORER_REGISTRY) == 7	P0
Aggregator — output clamp	Integration	Construct row that would produce weighted_sum > 1.0	conviction_signal ≤ 1.0	P0
Aggregator — per-scorer overflow guard	Integration	Individual scorer returns value > 1.0	Clamped before contributing to weighted_sum	P1
Cluster scorer — same-day discount	Unit	2 insiders same day vs 2 insiders different days	Different cluster_count, different score	P1
Cluster scorer — BUY vs SELL direction	Unit	Cluster of buyers vs cluster of sellers	Correct directional signal, not always bullish	P1
10b5-1 extraction — XML tag	Unit	Form 4 XML with <aff10b5One>1</aff10b5One>	is_10b51_plan = True	P0
10b5-1 extraction — footnote regex	Unit	Form 4 with "pursuant to a Rule 10b5-1" footnote	is_10b51_plan = True	P0

Test Case	Area	What to Test	Expected Result	Priority
10b5-1 — no flag present	Unit	Form 4 with no 10b5-1 indicators	is_10b51_plan = False	P0
Entity filer gate — parser tagging	Unit	Form 4 with <isCompany>1</isCompany>	is_company = True	P0
Entity filer gate — aggregator block	Integration	Row with is_company = True passed to aggregator	conviction_signal = 0.0, not scored	P0
Enrichment — arithmetic fallback	Unit	Row with shares_after only (no shares_before)	pct_holdings_transacted computed via fallback strategy	P1
Enrichment — forward-carry	Unit	Consecutive rows with same insider, missing shares_before	Previous row's shares_after carried forward	P1
YAML config — missing scorer	Config	Remove a scorer from YAML; run aggregator	UserWarning emitted, scorer skipped — no crash	P2
CSV output validation	Pipeline	Run full Phase 2 pipeline end-to-end	insider_signals_phase2.csv: 8,013 rows, no NaN in conviction_signal, all values in [−1.0, +1.0]	P0
Phase 1 pipeline re-run	Pipeline	Run legacy/run_mirror.py end-to-end	dissonance_scores.csv regenerated; 10 rows; scores 0–100	P2
Before/after signal comparison	Regression	Compare Phase 1 vs Phase 2 signal for same tickers	Phase 2 signals more granular; same directional bias	P2
10b5-1 penalty effect	Data Quality	Filter is_10b51_plan=True vs False, compare mean signal magnitude	True rows have lower |conviction_signal|	P1
Entity filer in production data	Data Quality	Filter known ETF/fund tickers in output CSV	Signal = 0.0 or not present	P1

SECTION 12

Recommended Next Steps

#	Action	Command / Detail	Priority
1	**Commit Phase 2 to git**	`git add src/ config/ tests/ docs/ legacy/ && git commit -m "Phase 2: conviction engine"`	Critical
2	**Verify .env is gitignored**	`cat .gitignore | grep .env` — confirm .env and *.env are listed	Critical
3	**Run full pytest suite**	`pytest tests/ -v` — confirm 31/31 passing	High
4	**Run Phase 2 pipeline end-to-end**	`python src/pipeline/run_scoring_pipeline.py`	High
5	**Validate output CSV**	Check: 8,013 rows, all conviction_signal in [−1.0, +1.0], no NaN	High
6	**Add deprecated header to root insider_parser.py**	Add `# DEPRECATED: use src/parsers/insider_parser_v2.py` at top of file	Medium
7	**Update README quickstart**	Point to `src/pipeline/run_scoring_pipeline.py` as Phase 2 entry point	Medium
8	**Run 10b5-1 penalty validation**	Filter output CSV for `is_10b51_plan=True` ; confirm lower mean |signal|	Medium
9	**Add run_timestamp to output CSV**	Prevents yfinance market data staleness ambiguity across multiple runs	Medium
10	**Implement technical_score scorer**	Price structure signal (close > 200DMA, 5–20% below 52w high). Enabled in YAML pending backtest.	Low
11	**Phase 3 — SQLite migration**	Replace flat CSV with SQLite for multi-run history, faster queries, multi-user access	Low

SECTION 13

Interview Explanation

The following is a 60–90 second spoken explanation suitable for a technical interview or mentor review session.

"When I built MIRROR, the original insider signal had a fundamental limitation:" *it mapped every trade to just five values based on raw dollar amounts — ±0.7, ±0.3, or zero. That means a CEO selling 80% of their entire stake got the same signal as someone selling a tiny fraction. There was no context.*

So in Phase 2, I rebuilt the scoring layer from scratch. *I added three enrichment stages: one that*

computes what percentage of the insider's holdings they traded, one that fetches market cap and trading volume from Yahoo Finance to normalize the trade size, and one that detects whether the trade was a pre-planned 10b5-1 program — which carries zero predictive value and should be penalized.

I then implemented a scorer registry pattern — each scorer is a decorated function that self-registers at import time. The aggregator pulls from that registry and applies weights from a YAML config file, so analysts can tune the model without touching Python. The output is a continuous conviction signal clamped to [−1.0, +1.0] instead of five discrete values.

The result: *8,013 transaction-level signals across 111 small/mid-cap tickers, with 1,533 10b5-1 trades correctly penalized and 31 tests confirming correctness. I also found and fixed a bug in the entity filer detection — the original code used a broken XPath and name heuristics when the XML tag was sitting right there. That fix prevents institutional non-discretionary trades from polluting the conviction signals.*

The architecture is now layered, testable, and extensible. *Adding a new scorer means creating one file with a decorator — nothing else changes.*

SECTION 14

Appendix

A. Key Output Files

File	Rows	Key Columns	Pipeline	Git
`results/insider_signals_phase2.csv`	8,013	ticker, insider_name, title, transaction_date, dollar_value, pct_holdings_transacted, market_cap, is_10b51_plan, conviction_signal	Phase 2	gitignored
`results/dissonance_scores.csv`	10	ticker, insider_signal, media_signal, options_signal, dissonance_score, risk_level	Phase 1	gitignored
`data/insider_transactions.csv`	8,013	Parsed Form 4 fields — shared input for both pipelines	Both	gitignored
`data/insider_filings.csv`	Needs verification	Filing index from SEC EDGAR for 146-ticker universe	Phase 2	gitignored

B. Key Configuration Files

File	Purpose	Format
`config/scoring.yaml`	Scorer weights and enable/disable flags — single source of truth	YAML
`config/universe_smallmid.txt`	146-ticker curated small/mid-cap universe for Phase 2 data pull	Plain text, one ticker per line
`.env`	API keys: NEWS_API_KEY (NewsAPI), SEC_API_KEY (if required)	KEY=VALUE, gitignored
`requirements.txt`	Python dependencies: numpy, pandas, requests, textblob, yfinance, PyYAML, pytest	pip format

C. Glossary

Term	Definition
Form 4	SEC regulatory filing required within 2 business days of any insider transaction. Contains insider name, role, number of shares, price, buy/sell direction, and 10b5-1 plan flag.

Term	Definition
10b5-1 plan	A pre-arranged trading plan that allows corporate insiders to trade on a fixed schedule, regardless of material non-public information. Carries zero informational content for conviction scoring.
Conviction signal	MIRROR's Phase 2 output — a float in [−1.0, +1.0] representing the strength and direction of an insider's informational signal. Positive = bullish; negative = bearish.
Dissonance score	MIRROR's Phase 1 output — a float 0–100 measuring contradiction between insider activity, news sentiment, and options positioning.
ADV	Average Daily Volume (dollar-denominated). Used to normalize trade size by liquidity.
Entity filer	An institutional filer (index fund, ESPP trust, custodian) that files Form 4 but makes non-discretionary trades. MIRROR hard-gates these to conviction_signal = 0.0.
Registry pattern	A design pattern where components self-register at import time via a decorator. MIRROR's @register_scorer decorator adds each scorer function to SCORER_REGISTRY when its module is imported.
pct_holdings_transacted	The percentage of an insider's total holdings in a company that were traded in a single transaction. The dominant conviction signal feature (weight: 0.30).
CIK	Central Index Key — SEC's unique identifier for each company and individual filer. Required to query Form 4 filings from EDGAR.

D. Scoring YAML — Annotated

```yaml
conviction_score:
  weight: 0.30     # Dominant scorer — % of holdings transacted
  enabled: true

role_score:
  weight: 0.20     # CEO=1.0, CFO=0.90, Director=0.55
  enabled: true

market_cap_score:
  weight: 0.15     # Trade size / market cap
  enabled: true

cluster_score:
  weight: 0.15     # Multi-insider confirmation
  enabled: true

liquidity_score:
  weight: 0.10     # Trade size / 30d ADV
  enabled: true

technical_score:
  weight: 0.10     # Price structure — DISABLED until backtested
  enabled: false

routine_penalty:
  weight: -0.20    # Calendar seasonality — zero predictive power
  enabled: true

tenb_penalty:
  weight: -0.25    # Pre-planned 10b5-1 trades — zero information
  enabled: true
```

```
# Net positive weight capacity: 0.90
# Net penalty capacity:        -0.45
# Effective range after clamp: [-1.0, +1.0]
```

E. is_company Bug — Code Before and After

BEFORE (BROKEN)

```
owner_name = (root.findtext(
  './/reportingOwner/'
  'reportingOwnerId/'   # ← node doesn't exist
in EDGAR XML
  'rptOwnerName') or '').upper()
name_words = set(...)
flags['is_company'] = bool(
  name_words & _ENTITY_KEYWORDS)
# Entity filers were being scored
# instead of gated to 0.0
```

AFTER (FIXED)

```
is_co = root.find(
  './/reportingOwner/'
  'reportingOwnerType/isCompany')
if is_co is not None and is_co.text:
    flags['is_company'] = (
      is_co.text.strip()
      in ('1', 'true', 'yes'))
else:
    # Fallback: name heuristics
    # for pre-2023 filings
    ...
```

PDF EXPORT INSTRUCTIONS

Open this file in **Google Chrome** or **Microsoft Edge**.
Press **Ctrl+P** (Windows) or **Cmd+P** (Mac).
Select **"Save as PDF"** as the destination.
Under "More settings", enable **"Background graphics"**.
Set margins to **None** for best results. Click Save.